

# CSC 411 Lecture 08: Generative Models for Classification

Ethan Fetaya, James Lucas and Emad Andrews


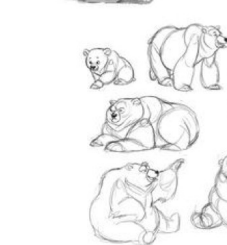
University of Toronto

# Today

- Classification – Bayes classifier
- Estimate **input** probability densities from data
- Naive Bayes


# Classification

- Given inputs  $\mathbf{x}$  and classes  $y$  we can do classification in several ways. How?

	(features)	(class label)
	$\mathbf{x}$  e.g: <ul style="list-style-type: none"><li>• height</li><li>• weight</li><li>• color</li></ul>	$y$  <b>elefant</b>
		<b>bear</b>

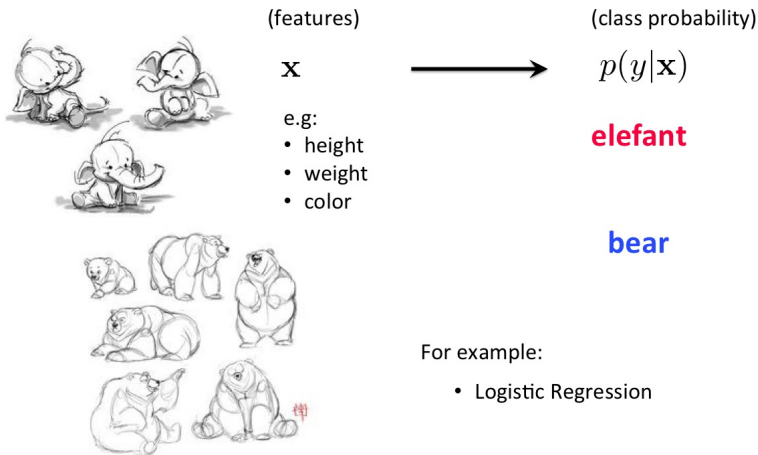
# Discriminative Classifiers

- **Discriminative** classifiers try to either:
  - ▶ learn mappings directly from the space of inputs  $\mathcal{X}$  to class labels  $\{0, 1, 2, \dots, K\}$



# Discriminative Classifiers

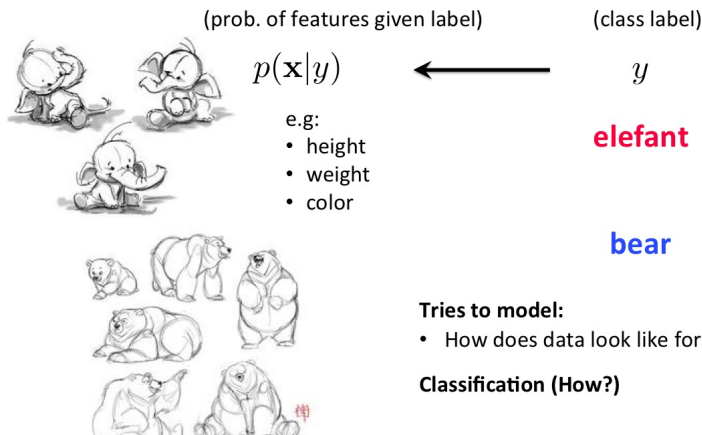
- **Discriminative** classifiers try to either:
  - ▶ or try to learn  $p(y|\mathbf{x})$  directly



# Generative Classifiers

How about this approach: build a model of “what data for a class looks like”

- **Generative** classifiers try to model  $p(\mathbf{x}, y)$ . If we know  $p(y)$  we can easily compute  $p(\mathbf{x}|y)$ .
- Classification via Bayes rule (thus also called Bayes classifiers)



# Generative vs Discriminative

Two approaches to classification:

- Discriminative classifiers estimate parameters of decision boundary/class separator directly from labeled examples. Tries to solve: How do I separate the classes?
  - ▶ learn  $p(y|\mathbf{x})$  directly (logistic regression models)
  - ▶ learn mappings from inputs to classes (least-squares, decision trees)
- Generative approach: model the distribution of inputs characteristic of the class (Bayes classifier). Tries to solve: What does each class "look" like?
  - ▶ Build a model of  $p(\mathbf{x}|y)$
  - ▶ Apply Bayes Rule

# Bayes Classifier

- Aim to classify text into spam/not-spam (yes  $C=1$ ; no  $C=0$ )
- Use bag-of-words features, get binary vector  $\mathbf{x}$  for each patient
- Given features  $\mathbf{x} = [x_1, x_2, \dots, x_d]^T$  we want to compute class probabilities using Bayes Rule:

$$p(C|\mathbf{x}) = \frac{p(\mathbf{x}|C)p(C)}{p(\mathbf{x})}$$

- More formally

$$\text{posterior} = \frac{\text{Class likelihood} \times \text{prior}}{\text{Evidence}}$$

- How can we compute  $p(\mathbf{x})$  for the two class case? (Do we need to?)

$$p(\mathbf{x}) = p(\mathbf{x}|C=0)p(C=0) + p(\mathbf{x}|C=1)p(C=1)$$

- To compute  $p(C|\mathbf{x})$  we need:  $p(\mathbf{x}|C)$  and  $p(C)$



# Classification: Simple Example

- Let's start with a simple (but slightly redundant) example.
- Imagine that we have some biased coins and we observe a single outcome from one of these coins.
- We have  $P(x|C) = \text{Ber}(\theta_C) = \theta_C^x \cdot (1 - \theta_C)^{1-x}$
- Notice that we have different parameters for each coin
- How can I fit the distribution to my data?
- Simple approach - maximum likelihood

# MLE for Bernoulli

- Assumption: data points are **independent** and **identically** distributed (i.i.d)

$$p(\mathcal{D}_C|C) = \prod_{n=1}^N p(x^{(n)}|C) = \prod_{n=1}^N \theta_C^{x^{(n)}} \cdot (1 - \theta_C)^{1-x^{(n)}} = \theta_C^{N_C} \cdot (1 - \theta_C)^{N-N_C}$$

- We define  $N_C = \sum_{i=1}^N x^{(i)}$  the number of ones (heads) seen.
- $N$  and  $N_C$  are called sufficient statistics - hold all the information we need to compute  $P(\mathcal{D}_C|C)$
- We can minimize the **negative log-likelihood** (NLL)





$$\ell_{\log-loss} = -\log(p(x^{(1)}, \dots, x^{(N)}|C)) = -N_C \log(\theta_C) - (N - N_C) \log(1 - \theta_C)$$

$$\frac{\partial \ell_{\log-loss}}{\partial \theta_C} = -\frac{N_C}{\theta_C} + \frac{N - N_C}{1 - \theta_C} = 0 \Rightarrow \theta_C = \frac{N_C}{N}$$

- MLE solution  $\theta_C = \frac{N_C}{N}$ . What if  $N_C = 0$ ?
- Example: Some rare word unseen in a training corpus.
- In that case  $P(x|C) = 0$  no matter what other information we have!
- Solution: A prior over  $\theta$ .
- Simple (conjugate) prior: Beta distribution
  - ▶  $\text{Beta}(\theta|a, b) \propto \theta^{a-1}(1 - \theta)^{b-1}$

# Beta Distribution

- Examples of  $Beta(\theta|a, b) \propto \theta^{a-1}(1 - \theta)^{b-1}$ :



- Likelihood  $p(\mathcal{D}_C|\theta_C) = \theta_C^{N_C} \cdot (1 - \theta_C)^{N - N_C}$
- Prior  $P(\theta_C) = \text{Beta}(\theta_C|a, b) \propto \theta_C^{a-1}(1 - \theta_C)^{b-1}$

$$\begin{aligned} p(\theta_C|\mathcal{D}_C) &= \frac{p(\mathcal{D}_C|\theta_C)P(\theta_C)}{p(\mathcal{D}_C)} \propto \theta_C^{N_C} \cdot (1 - \theta_C)^{N - N_C} \theta_C^{a-1}(1 - \theta_C)^{b-1} \\ &= \theta_C^{N_C + a - 1} \cdot (1 - \theta_C)^{N - N_C + b - 1} \end{aligned}$$

- We have  $P(\theta_C|\mathcal{D}_C) = \text{Beta}(N_C + a, N - N_C + b)$
- MAP estimation  $\theta_{C, \text{map}} = \frac{N_C + a - 1}{N + a + b - 2}$  (show!)

## Beta-Binomial \*

- Can we do better than using the MAP estimator? A more Bayesian approach.
- We have  $P(\theta_C | \mathcal{D}_C) = \text{Beta}(N_C + a, N - N_C + b)$ , what is  $P(x = 1 | \mathcal{D}_C)$ ?

$$\begin{aligned} P(x = 1 | \mathcal{D}_C) &= \int_0^1 P(x = 1 | \theta_C) P(\theta_C | \mathcal{D}_C) \\ &= \int_0^1 \theta_C P(\theta_C | \mathcal{D}_C) = \mathbb{E}[\theta_C | \mathcal{D}_C] \end{aligned}$$

- Beta(a,b) has a closed form mean  $\frac{a}{a+b}$  (a bit of work to show) so  $\theta_C = P(x = 1 | \mathcal{D}_C) = \frac{N_C + a}{N + a + b}$
- Equivalent to **pseudo-counts**, adding  $a$  fictitious positive examples and  $b$  negative ones.

# Moving beyond coins

- In the real world we tend to have a vector of observations  $\mathbf{x} = [x_1, \dots, x_d]$ .
- Modelling  $p(\mathbf{x}, y)$  in this case is much more complex.

$$p(x_1, \dots, x_d, y) = p(x_1 | x_2, \dots, x_d, y) \cdots p(x_{d-1} | x_d, y) p(x_d, y)$$

- We need to make some assumptions!
- The Naive-Bayes Model is born from a particularly strong assumption.

# Naive-Bayes for Bernoulli variables

- Make the (naive) assumption - dimensions  $\mathbf{x} = [x_1, \dots, x_d]$  are independent given the class  $y$ .

$$P(\mathbf{x}|y = C, \theta_C) = \prod_{j=1}^d p(x_j|y = C, \theta_{jC}) = \prod_{j=1}^d \theta_{jC}^{x_j} (1 - \theta_{jC})^{(1-x_j)} =$$

$$\exp \left( \sum_{j=1}^d x_j \log(\theta_{jC}/(1 - \theta_{jC})) + \sum_{j=1}^d \log(1 - \theta_{jC}) \right) = \exp(\mathbf{w}_C^T \mathbf{x} + w_{0C})$$

- Define  $w_{Cj} = \log(\theta_{jC}/(1 - \theta_{jC}))$ ,  $w_{0C} = \sum_{j=1}^d \log(1 - \theta_{jC})$



# Naive-Bayes for Bernoulli variables

- How do we classify?

$$P(y = C|\mathbf{x}) \propto P(y = C)P(\mathbf{x}|y = C) = \exp(\mathbf{w}_C^T \mathbf{x} + b_C)$$

- $w_{Cj} = \log(\theta_{jC}/(1 - \theta_{jC}))$ ,  $b_C = w_{0C} + \log(P(y = C))$
- Linear classifier! Model is similar to logistic regression, but different optimization.
  - ▶ No gradients - just need to count! Really fast to train.
  - ▶ Doesn't take into account correlation between features.

# Example: 20newsgroups

Table: Top word per topic

Topic	Naive Bayes	Logistic regression
'alt.atheism',	don	enlightening
'comp.graphics',	thanks	needed
'comp.os.ms-windows.misc',	windows	windows
'comp.sys.ibm.pc.hardware',	thanks	disappointing
'comp.sys.mac.hardware',	mac	mac
'comp.windows.x',	window	xtvaappinitialize
'misc.forsale',	sale	semd
'rec.autos',	car	car
'rec.motorcycles',	bike	bike
'rec.sport.baseball',	year	950k
'rec.sport.hockey',	team	hockey
'sci.crypt',	key	encryption
'sci.electronics',	use	cci
'sci.med',	don	melittin
'sci.space',	space	launch
'soc.religion.christian',	god	satan
'talk.politics.guns',	people	gun
'talk.politics.mideast',	people	kidding
'talk.politics.misc',	people	paranoia
'talk.religion.misc'	people	compuserve

# Beyond Bernoulli

- We focused on binary features,  $x_i$ , but Naive bayes is more general.
- Discrete features - multinomial.
- Continuous features - Gaussian (or any other).
- No problem to mix (unlike logistic regression)!

- Learning parameters:
  - ▶ Estimate  $P(y = C)$ , e.g.  $P(y = C) = \frac{\# \text{ class } C}{\# \text{ data points}}$
  - ▶ For each class  $C$  and feature  $x_i$  estimate the distribution  $p(x_i|y = C)$
- At test time:
  - ▶ For each class compute  $S_C = \log(P(Y = C)) + \sum_{i=1}^d \log(p(x_i|y = C))$
  - ▶ Classify according to  $\max_C S_C$
- Probabilities:  $P(y = C|\mathbf{x}) = \frac{\exp(S_C)}{\sum_{i=1}^L \exp(S_i)}$

- Pros:

- ▶ Really fast to train (single pass through data!).
- ▶ Fast to test.
- ▶ Less over-fitting, sometimes better than logistic on small data sets
- ▶ Easy to add/remove classes
- ▶ Can handle partial data.

- Cons:

- ▶ When naive i.i.d assumption doesn't hold (almost always) - can perform much worse.