

CSC 411 Lecture 18: Kernels

Ethan Fetaya, James Lucas and Emad Andrews

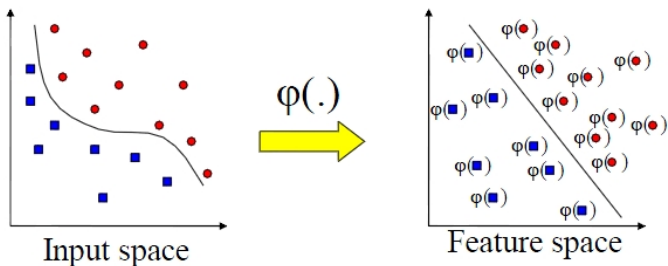
University of Toronto

Today

- Kernel trick
- Representer theorem

Non-linear decision boundaries

- We talk about SVM: **max margin linear classifier**
- Linear is limiting, how do we get non-linear decision boundaries?
- Feature mapping $\mathbf{x} \rightarrow \phi(\mathbf{x})$



- How do we find good features?
- If features are in a high dimension - high computational cost.

Motivation

- Let's say that we want a quadratic decision boundary
- What feature mapping do we need?
- One possibility (ignore arbitrary $\sqrt{2}$ for now)

$$\phi(\mathbf{x}) = (1, \sqrt{2}x_1, \dots, \sqrt{2}x_d, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \dots, \sqrt{2}x_{d-1}x_d, x_1^2, \dots, x_d^2)$$

Pairwise is over $i < j$

- We have $\dim(\phi(\mathbf{x})) = \mathcal{O}(d^2)$, could be problematic for large d .
- How can this be addressed?

Kernel Trick Idea

- Linear algorithms are based on inner-product
- What if you could compute the inner product without computing $\phi(\mathbf{x})$?
- Our previous example:

$$\phi(\mathbf{x}) = (1, \sqrt{2}x_1, \dots, \sqrt{2}x_d, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \dots, \sqrt{2}x_{d-1}x_d, x_1^2, \dots, x_d^2)$$

- What is $K(\mathbf{x}, \mathbf{y}) = \langle \phi(\mathbf{x}), \phi(\mathbf{y}) \rangle$?

$$\langle \phi(\mathbf{x}), \phi(\mathbf{y}) \rangle = 1 + \sum_{i=1}^d 2x_i y_i + \sum_{i,j=1}^d x_i x_j y_i y_j = (1 + \langle \mathbf{x}, \mathbf{y} \rangle)^2$$

- We can compute K in $\mathcal{O}(d)$ memory and compute time!
- K is called the (polynomial) [kernel](#).

Kernel SVM

- SVM dual form objective: $w = \sum \alpha_i t^{(i)} \mathbf{x}^{(i)}$

$$\max_{\alpha_i \geq 0} \left\{ \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i,j=1}^N t^{(i)} t^{(j)} \alpha_i \alpha_j (\mathbf{x}^{(i)T} \cdot \mathbf{x}^{(j)}) \right\}$$

$$\text{subject to } 0 \leq \alpha_i \leq C; \quad \sum_{i=1}^N \alpha_i t^{(i)} = 0$$

- Non-linear SVM using kernel function $K()$:

$$\max_{\alpha_i \geq 0} \left\{ \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i,j=1}^N t^{(i)} t^{(j)} \alpha_i \alpha_j K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \right\}$$

$$\text{subject to } 0 \leq \alpha_i \leq C; \quad \sum_{i=1}^N \alpha_i t^{(i)} = 0$$

- Unlike linear SVM, cannot express \mathbf{w} as linear combination of support vectors
 - ▶ now must retain the support vectors to classify new examples
- Final decision function:
$$y = \text{sign}[b + \langle \sum_{i=1}^N t^{(i)} \alpha_i \phi(\mathbf{x}^{(i)}), \phi(\mathbf{x}) \rangle] = \text{sign}[b + \sum_{i=1}^N t^{(i)} \alpha_i K(\mathbf{x}, \mathbf{x}^{(i)})]$$

- Examples of kernels: kernels measure similarity

1. Polynomial

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = (\mathbf{x}^{(i)T} \mathbf{x}^{(j)} + 1)^d$$

where d is the degree of the polynomial, e.g., $d = 2$ for quadratic

2. Gaussian/RBF

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2}{2\sigma^2}\right)$$

3. Sigmoid

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \tanh(\beta(\mathbf{x}^{(i)T} \mathbf{x}^{(j)})) + a)$$

- Kernel functions exist for non-vectorized data - string kernel, graph kernel, etc.
- Each kernel computation corresponds to a dot product
 - ▶ calculation for particular mapping $\phi(\mathbf{x})$ implicitly maps to high-dimensional space

Kernel Functions

- Mercer's Theorem (1909): any reasonable kernel corresponds to some feature space
- Reasonable means that the Gram matrix is positive semidefinite

$$K_{ij} = K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$$


- We can build complicated kernels so long as they are positive semidefinite.
- We can combine simple kernels together to make more complicated ones

Basic Kernel Properties

- Positive constant function is a kernel: for $\alpha \geq 0$, $K'(x_1, x_2) = \alpha$
- Positively weighted linear combinations of kernels are kernels: if $\forall i, \alpha_i \geq 0$,
 $K'(x_1, x_2) = \sum_i \alpha_i K_i(x_1, x_2)$
- Products of kernels are kernels: $K'(x_1, x_2) = K_1(x_1, x_2)K_2(x_1, x_2)$
- The above transformations preserve positive semidefinite functions
- We can use kernels as building blocks to construct complicated feature mappings

- Kernels let us express very large feature spaces
 - ▶ polynomial kernel $(1 + (\mathbf{x}^{(i)})^T \mathbf{x}^{(j)})^d$ corresponds to feature space exponential in d
 - ▶ Gaussian kernel has infinitely dimensional features
- Linear separators in these super high-dimensional spaces correspond to **highly non-linear decision boundaries** in the input space


Example - linear SVM



- Solid line - decision boundary. Dashed - $+1/-1$ margin. Purple - Bayes optimal
- Solid dots - Support vectors on margin

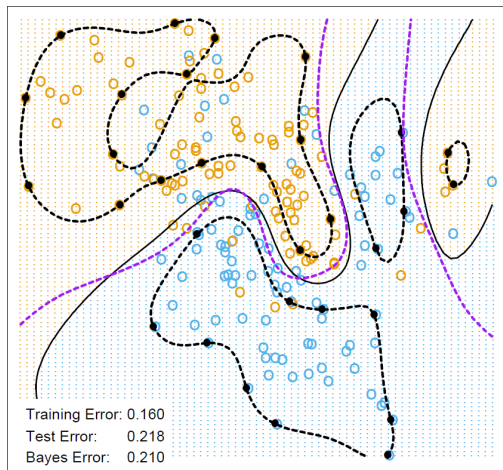
[Image credit: "Elements of statistical learning"]

Example - Deg 4 polynomial SVM



[Image credit: "Elements of statistical learning"]

Example - Gaussian SVM



[Image credit: "Elements of statistical learning"]

Kernel methods

- Kernels work well with SVM but not limited to it.
- When can we apply the kernel trick?

Representer Theorem:

If \mathbf{w}^* is defined as

$$\mathbf{w}^* = \arg \min \sum_{i=1}^N L(\langle \mathbf{w}, \phi(\mathbf{x}^{(i)}) \rangle, t^{(i)}) + \lambda \|\mathbf{w}\|^2$$

Then $\mathbf{w}^* \in \text{span}\{\phi(x_1), \dots, \phi(x_N)\}$, i.e. $\exists \alpha : \mathbf{w}^* = \sum_{i=1}^N \alpha_i \phi(x_i)$

- Proof idea: The subspace that is orthogonal to the span doesn't impact the loss, but increases the norm \Rightarrow Optimal thing is to set it to zero.
- We assume you can predict using inner-product.

- We can compute

$$\langle \mathbf{w}, \phi(\mathbf{x}) \rangle = \left\langle \sum_{i=1}^N \alpha_i \phi(\mathbf{x}^{(i)}), \phi(\mathbf{x}) \right\rangle = \sum_{i=1}^N \alpha_i \langle \phi(\mathbf{x}^{(i)}), \phi(\mathbf{x}) \rangle = \sum_{i=1}^N \alpha_i K(\mathbf{x}^{(i)}, \mathbf{x})$$

- Similarly for the regularizer

$$\begin{aligned} \|\mathbf{w}\|^2 &= \left\langle \sum_{i=1}^N \alpha_i \phi(\mathbf{x}^{(i)}), \sum_{j=1}^N \alpha_j \phi(\mathbf{x}^{(j)}) \right\rangle = \sum_{i,j=1}^N \alpha_i \alpha_j \langle \phi(\mathbf{x}^{(i)}), \phi(\mathbf{x}^{(j)}) \rangle \\ &= \sum_{i,j=1}^N \alpha_i \alpha_j K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \end{aligned}$$

- We can optimize without computing $\phi(\mathbf{x})$.

$$\alpha = \arg \min \sum_{i=1}^N L \left(\sum_{j=1}^N \alpha_j k(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}), t^{(i)} \right) + \lambda \sum_{i,j=1}^N \alpha_i \alpha_j K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$$

Other Kernel methods

- Kernel Logistic regression

- ▶ We can think of logistic regression as minimizing $\log(1 + \exp(-t^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}))$
- ▶ If you use L_2 regularization (Gaussian prior) this fits the representer theorem.
- ▶ Performance is close to SVM

- PCA

- ▶ A bit trickier to show how to only use kernels.
- ▶ Equivalent to first using a non-linear transformation to high dimension then use linear projection to low dimension.

- Kernel Bayesian methods (not covered in this course)

- ▶ Gaussian processes

Kernel and SVM

- The kernel trick is not limited to SVM, but is most common with it.
- Why do the kernel trick and SVM work well together?
- Generalization:
 - ▶ The kernel trick allows you to work in very high dimensions - what about overfitting?
 - ▶ SVM enjoys generalization bounds that don't depend on dimension (depend on margin or #support vectors).
 - ▶ Regularization is still very important to reduce overfitting.
- Computation:
 - ▶ In general \mathbf{w}^* is a linear combination of the training data
 - ▶ SVM only need to save a (hopefully small) subset of support vectors - Less memory and faster predictions.

- Advantages:

- ▶ Kernels allow very flexible hypotheses
- ▶ Kernel trick allows us to work in very high (or infinite) dimensional space
- ▶ Soft-margin extension permits mis-classified examples
- ▶ Can usually outperform linear svm

- Disadvantages:

- ▶ Must choose kernel parameters
- ▶ Large number of support vector \Rightarrow Computationally expensive to predict new points.
- ▶ Can overfit.

More Summary

- Software:
 - ▶ Sklearn implementation is based on LIBSVM (SMO algorithm)
 - ▶ SVMlight is among the earliest implementations
 - ▶ svm-Perf uses Cutting-Plane Subspace Pursuit.
 - ▶ Several Matlab toolboxes for SVM are also available
- Key points:
 - ▶ Difference between logistic regression and SVMs
 - ▶ Maximum margin principle
 - ▶ Target function for SVMs
 - ▶ Slack variables for mis-classified points
 - ▶ Kernel trick allows non-linear generalizations